

# Deep Reinforcement Learning for Process Control in Powder Bed Fusion

Supervisor

*George Panoutsos*

Candidate

*Ricardo Dominguez-Olmedo*



Department of Automatic Control & Systems Engineering  
The University of Sheffield

15th May 2019

# Table of contents

1. Background
2. Problem formulation
3. Results
4. Conclusions

## Background

---

# Deep reinforcement learning (DRL)

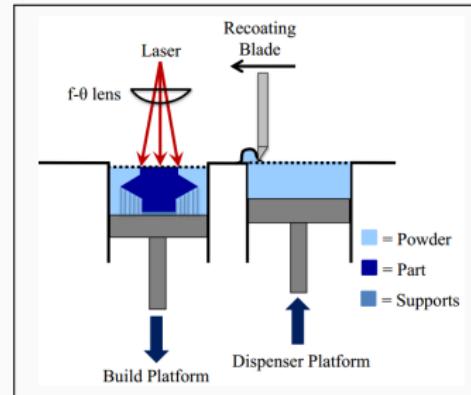
- Control learned autonomously through **interaction**
  - Observe consequence of actions and adapt behaviour
- No system identification or controller design required
  - Time consuming
  - Application-specific knowledge necessary
- Has only been used in robotic systems



Robotic grasping task

# Powder bed fusion (PBF)

- Additive manufacturing technology
  - 3D objects build by stacking layers of material
  - Very high geometric freedom
- Difficult to reliably manufacture high-quality products
  - “Optimal” parameters found by trial-and-error
  - Closed-loop control necessary to improve quality (surface temperature?)



Powder bed fusion process



Radial compressor

# Project aim

Use DRL for thermal control of a PBF process.

- Assess suitability of DRL for industrial process control
- PBF is very challenging...
  - Highly non-linear
  - High-dimensional
  - No known model
  - Expensive

## Problem formulation

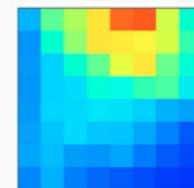
---

# Manufacturing system

- Controllable parameters: laser power and scan speed
- Pyrometer to measure surface temperature
  - Heat map of the entire layer surface
  - Divide into “cells” → **system state**
- **Objective:** maintain an optimal surface temperature



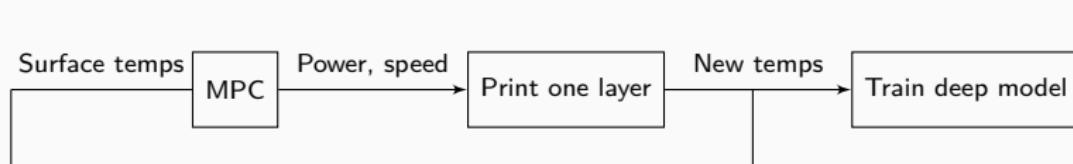
The AconityMINI machine



Surface heat map  
(red is hottest)

# The PETS Algorithm

- Iteratively train a model of the dynamics using data collected so far
  - Model is a **deep neural network**
- Use **Model Predictive Control** (MPC) over learned model
  - Control inputs over plan horizon optimised using a black-box optimiser
  - Cost function: exponential negative quadratic error  $-e^{-k(x-r)^T(x-r)}$



# Experiments

- Two different materials:
  - CM247, 10x10x5mm cuboids
  - Haynes 282, 5x5x5mm cuboids
- 20 cuboids built per experiment
  - 10 with constant optimal parameters
  - 10 with thermal control

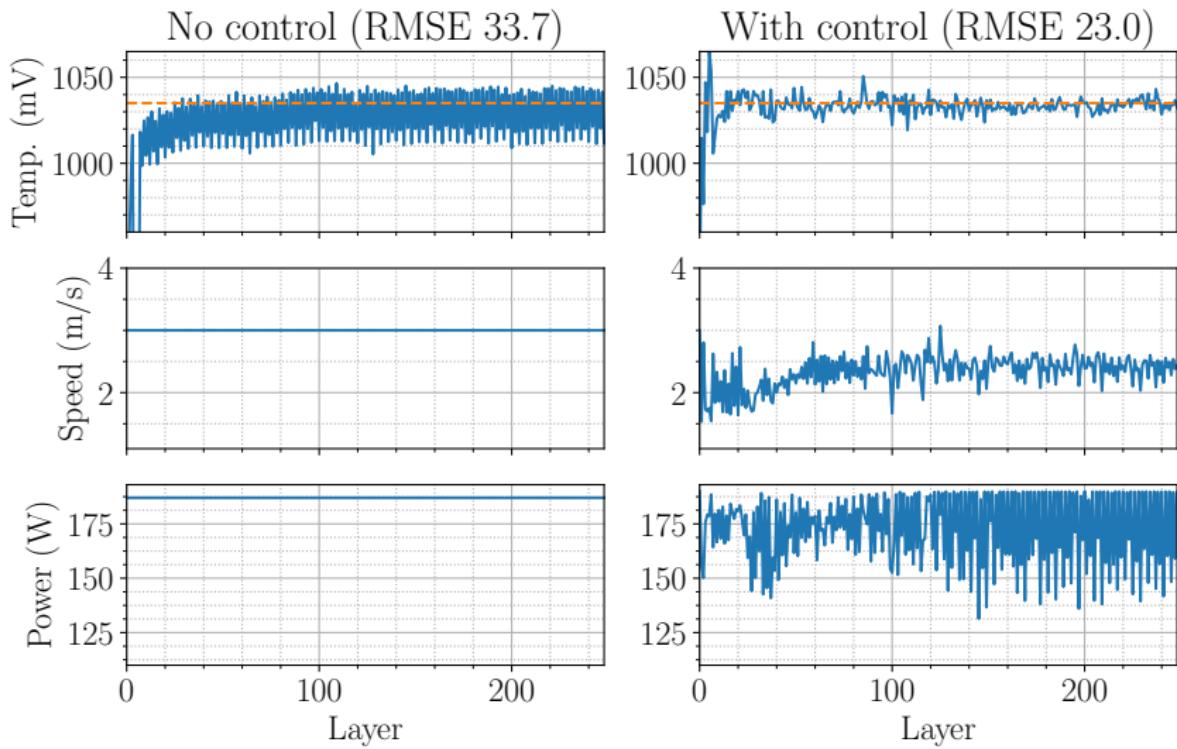


Sample cuboid

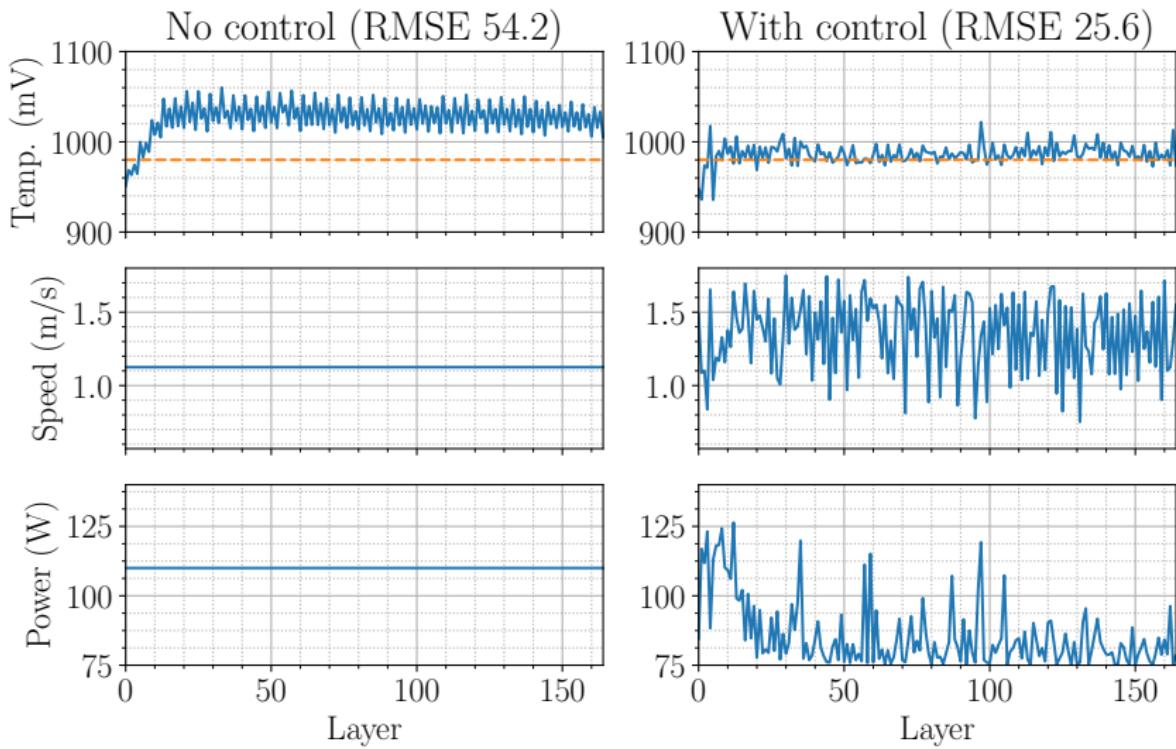
## Results

---

# Experiment I: CM247

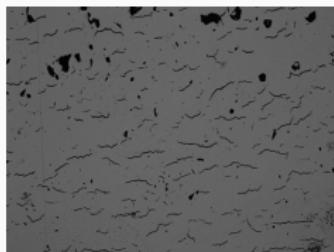


## Experiment II: Haynes 282

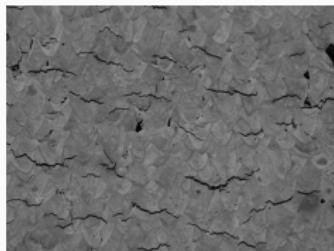


# Build quality

CM247



No control



With control

Haynes 282

Density of the samples built.

	Density (%)	
	Mean	Minimum
No control	<b><math>97.9 \pm 0.6</math></b>	97.2
With control	$95.9 \pm 0.7$	95.4

Why does it not improve?

- Relationship temperature - quality not well understood
- Drastic changes of input parameters

## Conclusions

---

## Advantages of our DRL approach in PBF

- Determining optimal parameters via trial-and-error is expensive
  - Tens of thousands of pounds and days
  - Redone for every new object manufactured
- DRL leads to good quality at much lower cost
- Quality - Cost trade-off very appealing in some instances

## Novel contributions

- 1<sup>st</sup> study to use DRL for process control of a real-world system
  - Very strong control performance for a complex process
  - No system identification or controller tuning required
- 1<sup>st</sup> study to demonstrate closed-loop thermal control in PBF
  - Research: further study relationship thermal signatures - build quality
  - Industry: acceptable product quality without expensive parameter tuning

Questions?

## PETS Algorithm: General

- 1: Initialise dataset  $\mathbb{D}$  (e.g. data from a single trial using a random control inputs)
- 2: **for** Trial  $k = 1$  to  $K$  **do**
- 3:     Train dynamics model  $\tilde{f}$  given  $\mathbb{D}$ .
- 4:     **for** Time  $t = 0$  to task time-horizon  $H$  **do**
- 5:         Obtain optimal control inputs  $u_{t:t+T}^*$  using CEM and TS
- 6:         Execute first control input  $u_t^*$
- 7:         Store outcome:  $\mathbb{D} \leftarrow \mathbb{D} \cup \{x_t, u_t^*, x_{t+1}\}$

$$\tilde{f}(x_{t+1} | x_t, u_t) = Pr(x_{t+1} | x_t, u_t) = \mathcal{N}(\mu(x_t, u_t), \sigma(x_t, u_t)) \quad (1)$$

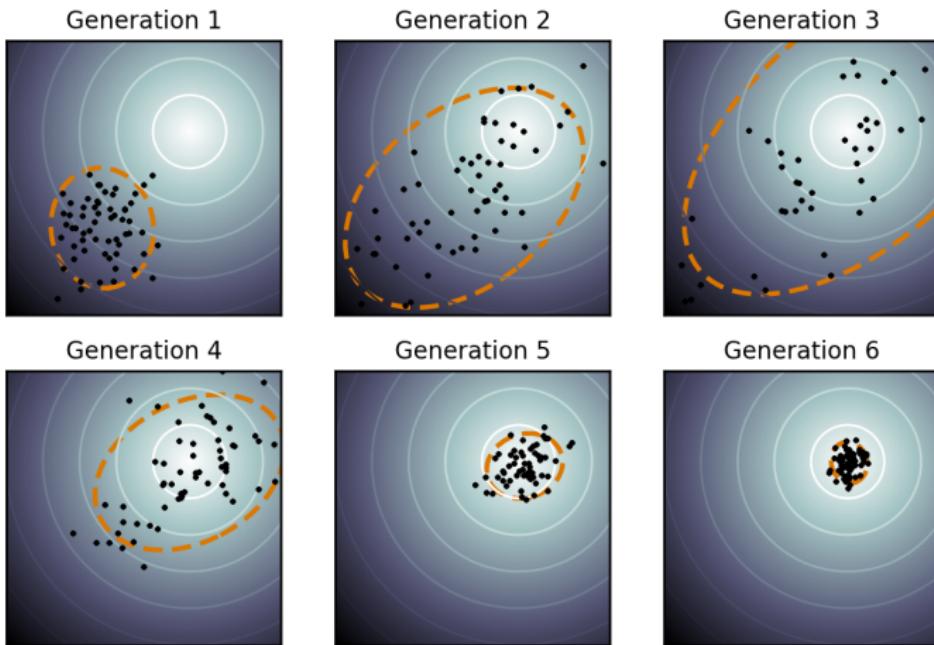
## PETS Algorithm: Trajectory Sampling

```
1: for Particle  $p = 1$  to  $P$  do
2:   Initialise particle to current system state  $x_t^p \leftarrow x_t$ 
3:   Uniformly sample a bootstrap  $b \leftarrow \{1, \dots, B\}$ 
4:   for Time-step  $\tau = t$  to  $t + T$  do
5:     Sample next system state  $x_{\tau+1}^p \leftarrow \tilde{f}_{\theta_b}(x_\tau^p, u_\tau)$ 
6:     if TS variant is TS1 then
7:       Uniformly resample a bootstrap  $b \leftarrow \{1, \dots, B\}$ 
8: return Expected cost  $J \leftarrow -\frac{1}{P} \sum_{p=0}^P \sum_{\tau=t}^{t+T} \mathcal{R}(x_{\tau+1}^p, u_\tau)$ 
```

# PETS Algorithm: Optimisation

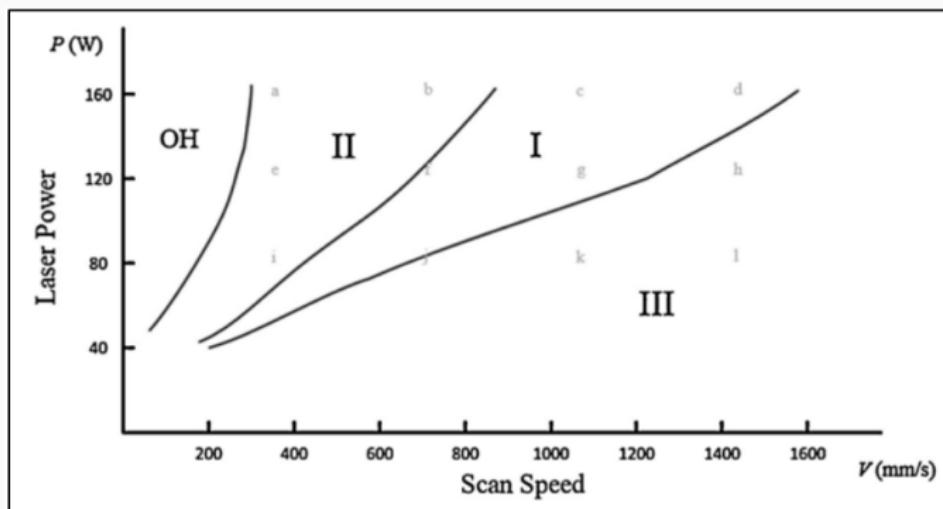
- 1: **for**  $i = 1$  to MaxIters **do**
- 2:     Sample  $N$  control input sequences  
 $u_{t:t+T}^1, \dots, u_{t:t+T}^N \sim \mathcal{N}(\mu_{u_{t:t+T}}, \sigma_{u_{t:t+T}}^2)$
- 3:     Evaluate expected cost  $J_1, \dots, J_N$  of  $u_{t:t+T}^1, \dots, u_{t:t+T}^N$  using TS
- 4:     Pick the  $N_{\text{elites}}$  control input sequences  $u'_{t:t+T}^1, \dots, u'_{t:t+T}^{N_{\text{elites}}}$  with  
lowest cost
- 5:     Compute mean  $\mu'_{u_{t:t+T}}$  and variance  $\sigma'^2_{u_{t:t+T}}$  of elites  $u'_{t:t+T}^1, \dots, u'_{t:t+T}^{N_{\text{elites}}}$ .
- 6:     Update mean  $\mu_{u_{t:t+T}} \leftarrow \alpha \mu_{u_{t:t+T}} + (1 - \alpha) \mu'_{u_{t:t+T}}$
- 7:     Update variance  $\sigma^2_{u_{t:t+T}} \leftarrow \alpha \sigma^2_{u_{t:t+T}} + (1 - \alpha) \sigma'^2_{u_{t:t+T}}$
- 8:     Constrain variance  $\sigma^2_{u_{t:t+T}} \leftarrow \min(\sigma^2_{u_{t:t+T}}, \frac{(\mu_{u_{t:t+T}} - u_{\min})^2}{4}, \frac{(u_{\max} - \mu_{u_{t:t+T}})^2}{4})$
- 9:     **if** Convergence criteria is met:  $\max(\sigma^2_{u_{t:t+T}}) \leq \epsilon$  **then**
- 10:       **break**
- 11: **return** Optimal control input sequence  $u_{t:t+T}^* \leftarrow \mu_{u_{t:t+T}}$

# CEM optimisation



CEM Optimisation

# Optimal energy envelopes in PBF



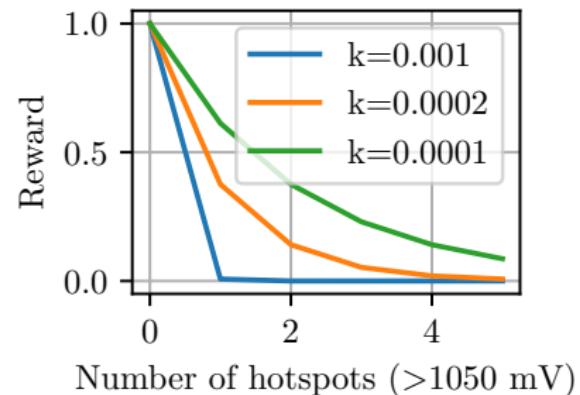
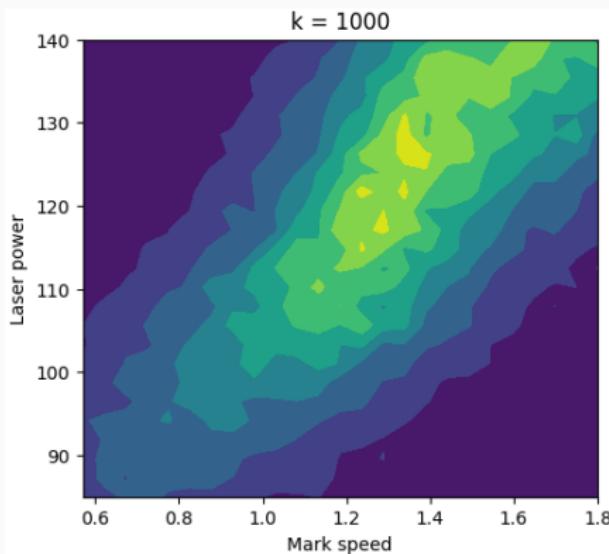
Example of energy density envelopes [?].

# Choice of number of cells

Model accuracy and number of cells

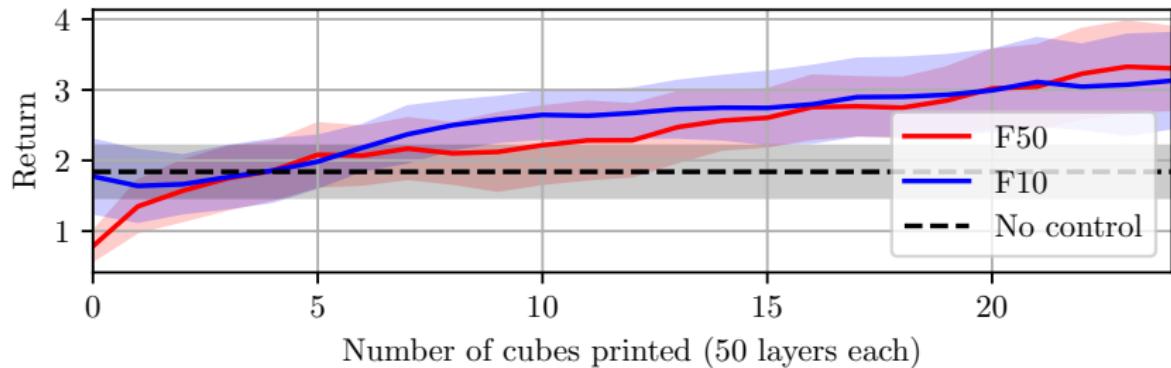
Number of cells	MSE
4	389±12
16	<b>274±11</b>
64	293±7

## Choice of $k$



- (a) Energy density envelopes for  $k = 0.001$ . Darker colours indicate lower return.  
(b) Reward for different numbers of hotspots.

## PETS sample-efficiency tests in simulation



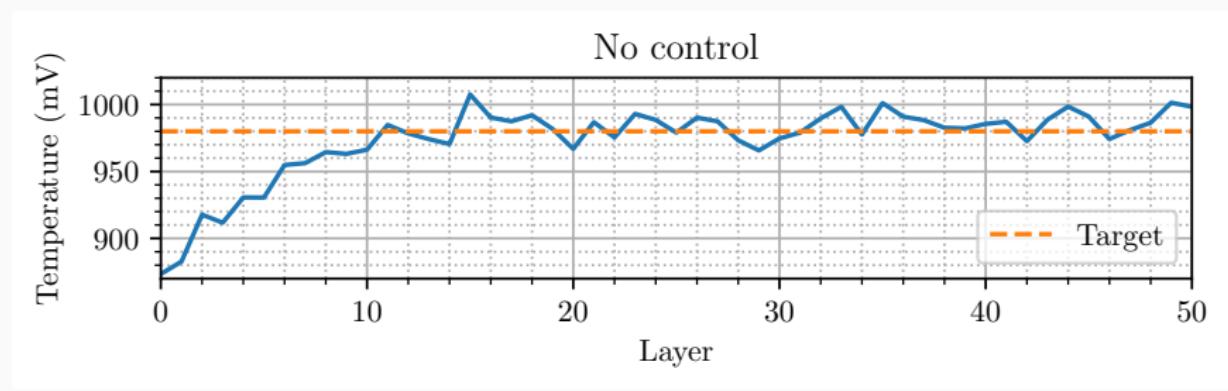
PETS control performance with respect to the number of cubes printed. The solid lines represent the mean return obtained for the 10 runs performed, while the shaded areas represent one standard deviation. "No control" refers to the use of constant optimal parameters.

# Controller performance comparison in simulation

Metrics for the Haynes 282 experiment (over 20 trials).

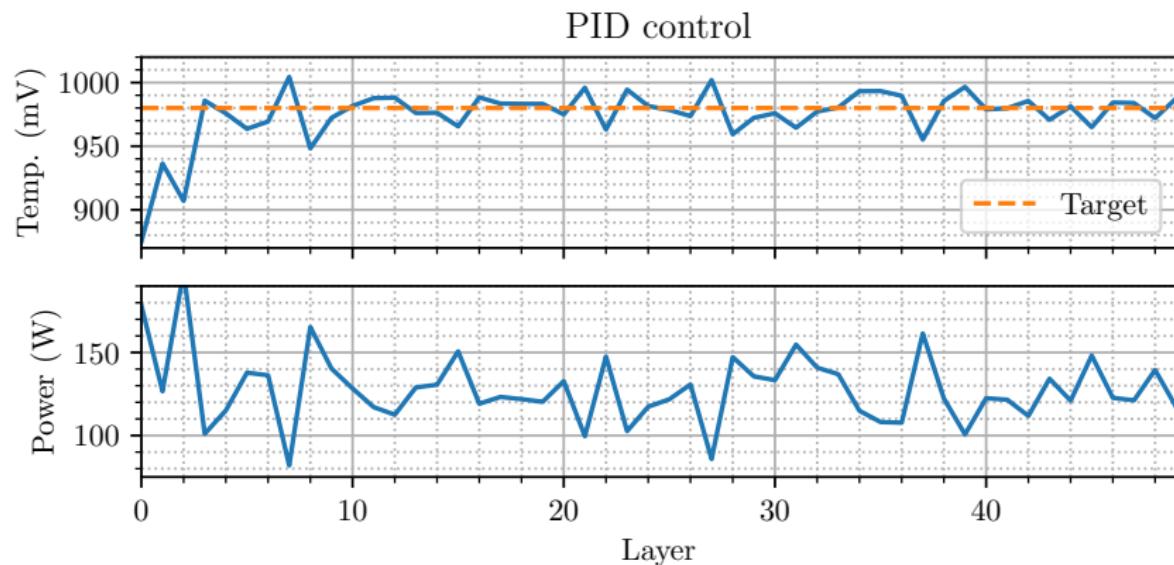
	RMSE	> +50 (%)	StdDev	Return
No control	$35.1 \pm 3.6$	$1.0 \pm 0.4$	$12.6 \pm 0.4$	$1.21 \pm 0.56$
PI control	$26.6 \pm 1.1$	$1.0 \pm 0.6$	$14.2 \pm 0.5$	$1.55 \pm 0.27$
PETS control	$24.9 \pm 0.6$	$0.1 \pm 0.1$	$12.0 \pm 0.3$	$4.71 \pm 0.92$

## Performance of no control in simulation



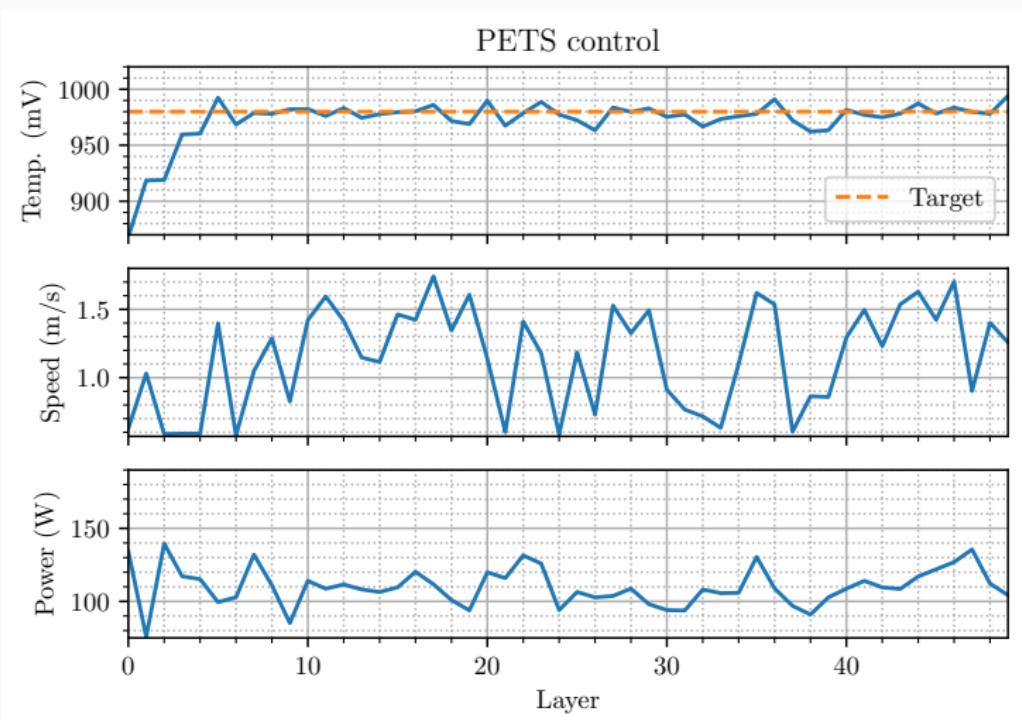
Time response of fixed optimal process parameters.

## Performance PID simulation

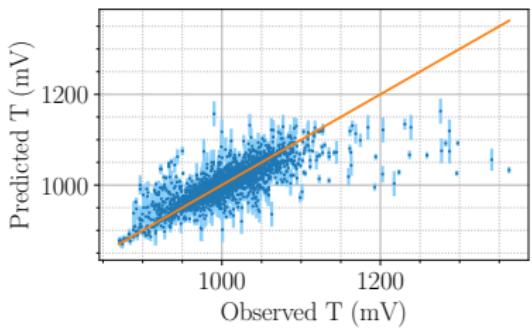


Time response for PID control.

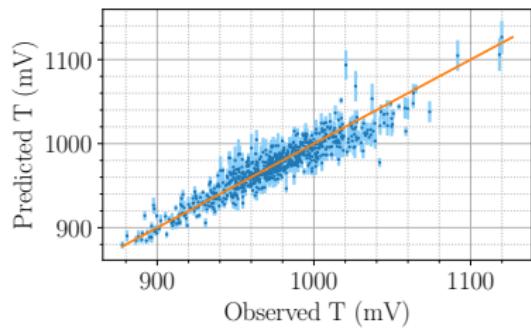
# Performance PETS simulation



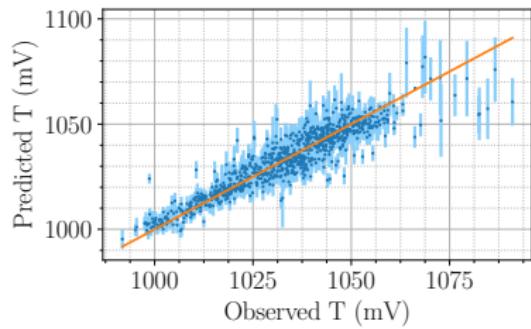
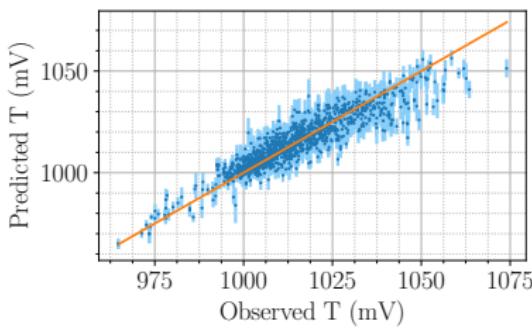
# Modelling Plots



Build 1.1



Build 2.1



# Modelling Table

Model performance for the different builds performed.

Build	Datapoints	MSE	R <sup>2</sup>
1.1	13972	844±94	0.627±0.021
1.2	4776	621±143	0.675±0.029
2.1	3948	204±38	0.834±0.038
2.2	8780	232±57	0.807±0.035
3	6280	50±5	0.849±0.024
4	4820	54±5	0.840±0.028